

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

In [2]:

```
from sklearn.datasets import load_boston
boston_datasets=load_boston()
```

In [4]:

```
print(boston_datasets.keys())
```

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

In [5]:

```
boston_datasets.DESCR
```

Out[5]:

```
".. _boston_dataset:\n\nBoston house prices dataset\n-----
-----\n\n**Data Set Characteristics:** \n\n      :Number of Instances: 506
\n\n      :Number of Attributes: 13 numeric/categorical predictive. Median V
alue (attribute 14) is usually the target.\n\n      :Attribute Information
(in order):\n          - CRIM      per capita crime rate by town\n          - ZN      proportion of residential land zoned for lots over 25,000 sq.ft.\n          - INDUS    proportion of non-retail business acres per town\n          - CHAS     Charles River dummy variable (= 1 if tract bounds river; 0 otherwis
e)\n          - NOX     nitric oxides concentration (parts per 10 million)
\n          - RM      average number of rooms per dwelling\n          - AGE      proportion of owner-occupied units built prior to 1940\n          - DIS      weighted distances to five Boston employment centres\n          - RAD      index of accessibility to radial highways\n          - TAX      full-value property-tax rate per $10,000\n          - PTRATIO  pupil-teacher ratio by town\n          - B        1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town\n          - LSTAT    % lower status of the population\n          - MEDV     Median value of owner-occupied homes in $1000's\n\n      :Missing Attribute Values: None\n\n      :Creator: Harrison, D. and Rubinfeld, D.L.\n\nThis is a copy of UCI ML housing dataset.\nhttps://archive.ics.uci.edu/ml/machine-learning-databases/housing/\n\n\nThis dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.\n\nThe Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol. 1.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics'\n...', Wiley, 1980. N.B. Various transformations are used in the table on\npages 244-261 of the latter.\n\nThe Boston house-price data has been used in many machine learning papers that address regression\nproblems.\n\n\n.. topic:: References\n\n    - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.\n    - Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.\n"
```


In [8]:

```
boston=pd.DataFrame(boston_datasets.data,columns=boston_datasets.feature_names)
boston
```

Out[8]:

|     | CRIM    | ZN   | INDUS | CHAS | NOX   | RM    | AGE  | DIS    | RAD | TAX   | PTRATIO | B      |
|-----|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|
| 0   | 0.00632 | 18.0 | 2.31  | 0.0  | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3    | 396.90 |
| 1   | 0.02731 | 0.0  | 7.07  | 0.0  | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8    | 396.90 |
| 2   | 0.02729 | 0.0  | 7.07  | 0.0  | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8    | 392.83 |
| 3   | 0.03237 | 0.0  | 2.18  | 0.0  | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7    | 394.63 |
| 4   | 0.06905 | 0.0  | 2.18  | 0.0  | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7    | 396.90 |
| ... | ...     | ...  | ...   | ...  | ...   | ...   | ...  | ...    | ... | ...   | ...     | ..     |
| 501 | 0.06263 | 0.0  | 11.93 | 0.0  | 0.573 | 6.593 | 69.1 | 2.4786 | 1.0 | 273.0 | 21.0    | 391.99 |
| 502 | 0.04527 | 0.0  | 11.93 | 0.0  | 0.573 | 6.120 | 76.7 | 2.2875 | 1.0 | 273.0 | 21.0    | 396.90 |
| 503 | 0.06076 | 0.0  | 11.93 | 0.0  | 0.573 | 6.976 | 91.0 | 2.1675 | 1.0 | 273.0 | 21.0    | 396.90 |
| 504 | 0.10959 | 0.0  | 11.93 | 0.0  | 0.573 | 6.794 | 89.3 | 2.3889 | 1.0 | 273.0 | 21.0    | 393.45 |
| 505 | 0.04741 | 0.0  | 11.93 | 0.0  | 0.573 | 6.030 | 80.8 | 2.5050 | 1.0 | 273.0 | 21.0    | 396.90 |

506 rows × 13 columns




In [11]:

```
boston['MEDV']=boston_datasets.target
boston
```

Out[11]:

|     | CRIM    | ZN   | INDUS | CHAS | NOX   | RM    | AGE  | DIS    | RAD | TAX   | PTRATIO | B      |
|-----|---------|------|-------|------|-------|-------|------|--------|-----|-------|---------|--------|
| 0   | 0.00632 | 18.0 | 2.31  | 0.0  | 0.538 | 6.575 | 65.2 | 4.0900 | 1.0 | 296.0 | 15.3    | 396.90 |
| 1   | 0.02731 | 0.0  | 7.07  | 0.0  | 0.469 | 6.421 | 78.9 | 4.9671 | 2.0 | 242.0 | 17.8    | 396.90 |
| 2   | 0.02729 | 0.0  | 7.07  | 0.0  | 0.469 | 7.185 | 61.1 | 4.9671 | 2.0 | 242.0 | 17.8    | 392.83 |
| 3   | 0.03237 | 0.0  | 2.18  | 0.0  | 0.458 | 6.998 | 45.8 | 6.0622 | 3.0 | 222.0 | 18.7    | 394.63 |
| 4   | 0.06905 | 0.0  | 2.18  | 0.0  | 0.458 | 7.147 | 54.2 | 6.0622 | 3.0 | 222.0 | 18.7    | 396.90 |
| ... | ...     | ...  | ...   | ...  | ...   | ...   | ...  | ...    | ... | ...   | ...     | ..     |
| 501 | 0.06263 | 0.0  | 11.93 | 0.0  | 0.573 | 6.593 | 69.1 | 2.4786 | 1.0 | 273.0 | 21.0    | 391.99 |
| 502 | 0.04527 | 0.0  | 11.93 | 0.0  | 0.573 | 6.120 | 76.7 | 2.2875 | 1.0 | 273.0 | 21.0    | 396.90 |
| 503 | 0.06076 | 0.0  | 11.93 | 0.0  | 0.573 | 6.976 | 91.0 | 2.1675 | 1.0 | 273.0 | 21.0    | 396.90 |
| 504 | 0.10959 | 0.0  | 11.93 | 0.0  | 0.573 | 6.794 | 89.3 | 2.3889 | 1.0 | 273.0 | 21.0    | 393.45 |
| 505 | 0.04741 | 0.0  | 11.93 | 0.0  | 0.573 | 6.030 | 80.8 | 2.5050 | 1.0 | 273.0 | 21.0    | 396.90 |

506 rows × 14 columns



In [12]:


```
boston.isnull().sum()
```

Out[12]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV      0
dtype: int64
```

In [14]:

```
sns.set(rc={'figure.figsize':(11.7,8.27)})
sns.distplot(boston['MEDV'],bins=30)
plt.show()
```




In [15]:

```
correlation_matrix=boston.corr().round(2)
sns.heatmap(data=correlation_matrix,annot=True)
```

Out[15]:

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x20b3046a408&gt;




In [16]:

```
plt.figure(figsize=(20,5))
features=['LSTAT','RM']
target=boston['MEDV']
```

&lt;Figure size 1440x360 with 0 Axes&gt;

In [17]:

```
for i,col in enumerate(features):  
    plt.subplot(1,len(features),i+1)  
    x=boston[col]  
    y=target  
    plt.scatter(x,y,marker='o')  
    plt.title(col)  
    plt.xlabel(col)  
    plt.ylabel('MEDV')
```



In [21]:

```
X=pd.DataFrame(np.c_[boston['LSTAT'],boston['RM']],columns=['LSTAT','RM'])  
Y=boston['MEDV']
```

In [22]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=5)
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

```
(404, 2)
(102, 2)
(404,)
(102,)
```

In [28]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
lin_model=LinearRegression()
lin_model.fit(X_train,Y_train)
```

Out[28]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

In [37]:

```
from sklearn.metrics import r2_score
y_train_predict=lin_model.predict(X_train)
rmse=(np.sqrt(mean_squared_error(Y_train,y_train_predict)))
r2=r2_score(Y_train,y_train_predict)
print("the model performance for training set")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")
y_test_predict=lin_model.predict(X_test)
rmse=(np.sqrt(mean_squared_error(Y_test,y_test_predict)))
r2=r2_score(Y_test,y_test_predict)
print("the model performance for testing set")
print('RMSE is {}'.format(rmse))
print('R2 score is {}'.format(r2))
print("\n")
```

```
the model performance for training set
RMSE is 5.6371293350711955
R2 score is 0.6300745149331701
```

```
the model performance for testing set
RMSE is 5.137400784702911
R2 score is 0.6628996975186953
```

In [ ]: